

TEMPORAL DATABASES: THEORY AND POSTGRES

Paul A. Jungwirth

31 May 2019

PGCon 2019

TEMPORAL IS DISTINCT FROM TIME-SERIES

time-series	temporal
single timestamp	two timestamps
records events	records things
IoT sensors, finance	auditing, history
challenge is scale	challenge is complexity
Partitioning	ranges, exclusion constraints
Citus, TimescaleDB	Teradata, temporal_tables

TEMPORAL = HISTORICAL

- e-commerce: product price

TEMPORAL = HISTORICAL

- e-commerce: product price
- real estate: house renovations

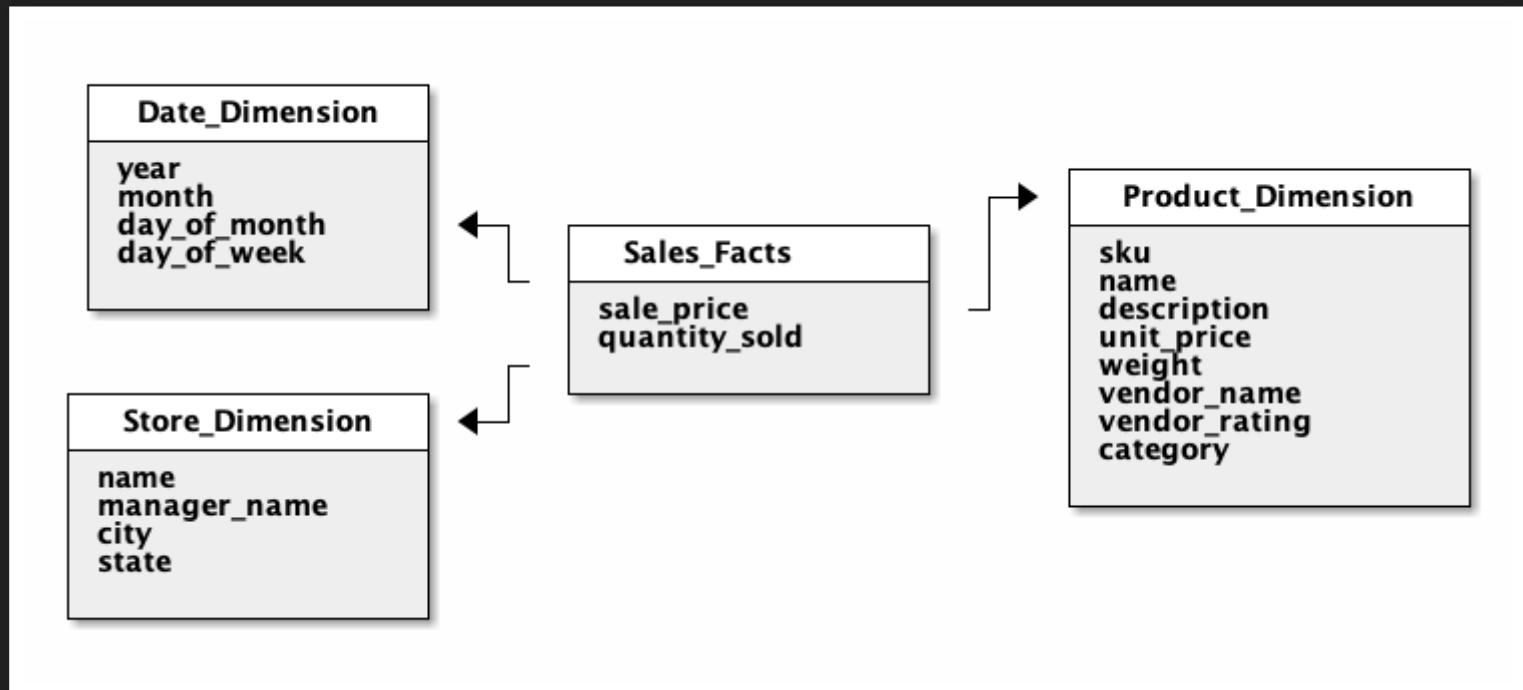
TEMPORAL = HISTORICAL

- e-commerce: product price
- real estate: house renovations
- employees: position, salary, employment period


TEMPORAL = HISTORICAL

- e-commerce: product price
- real estate: house renovations
- employees: position, salary, employment period
- questionnaires: changing questions, options

OLAP PROBLEMS TOO




OLAP PROBLEMS TOO



"SLOWLY-CHANGING DIMENSIONS"

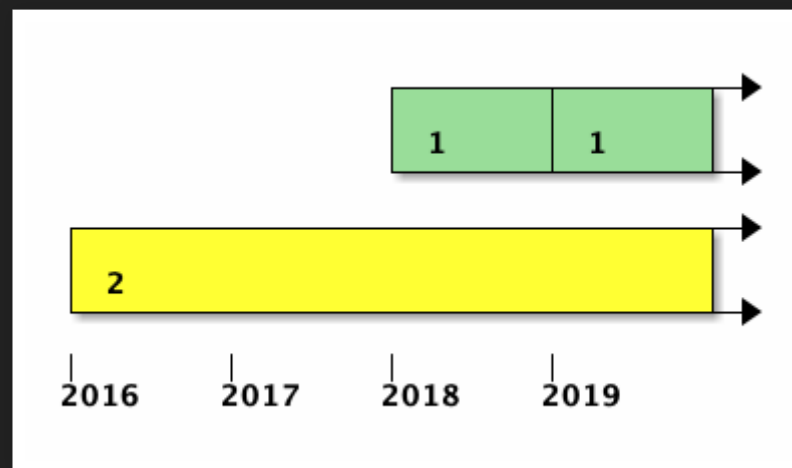
- Type I: Overwrite it
- Type II: Add a Row
- Type III: Add a Column

RESEARCH



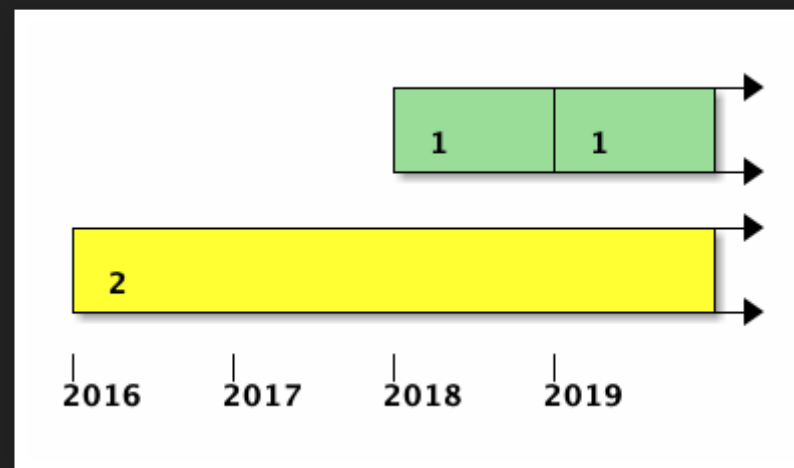
TEMPORAL EXAMPLE

products				
id	name	price	valid_from	valid_til
1	shoe	\$5	Jan 2018	Jan 2019
1	shoe	\$7	Jan 2019	
2	snow	\$2	Jan 2016	



TEMPORAL EXAMPLE

products			
id	name	price	valid_at
1	shoe	\$5	[Jan 2018, Jan 2019) [Jan 2019,)
1	shoe	\$7	
2	snow	\$2	[Jan 2016,)



RANGE OPERATORS

Operator	Description	Example	Result
=	equal	<code>int4range(1,5) = '[1,4]':int4range</code>	t
<>	not equal	<code>numrange(1.1,2.2) <> numrange(1.1,2.3)</code>	t
<	less than	<code>int4range(1,10) < int4range(2,3)</code>	t
>	greater than	<code>int4range(1,10) > int4range(1,5)</code>	t
<=	less than or equal	<code>numrange(1.1,2.2) <= numrange(1.1,2.2)</code>	t
>=	greater than or equal	<code>numrange(1.1,2.2) >= numrange(1.1,2.0)</code>	t

MORE OPERATORS

@>	contains range	int4range(2,4) @> int4range(2,3)	t
@>	contains element	'[2011-01-01,2011-03-01)':tsrange @> '2011-01-10':timestamp	t
<@	range is contained by	int4range(2,4) <@ int4range(1,7)	t
<@	element is contained by	42 <@ int4range(1,7)	f
&&	overlap (have points in common)	int8range(3,7) && int8range(4,12)	t
<<	strictly left of	int8range(1,10) << int8range(100,110)	t
>>	strictly right of	int8range(50,60) >> int8range(20,30)	t

AND MORE

&<	does not extend to the right of	<code>int8range(1,20) &< int8range(18,20)</code>	<code>t</code>
&>	does not extend to the left of	<code>int8range(7,20) &> int8range(5,10)</code>	<code>t</code>
- -	is adjacent to	<code>numrange(1.1,2.2) - - numrange(2.2,3.3)</code>	<code>t</code>
+	union	<code>numrange(5,15) + numrange(10,20)</code>	<code>[5,20)</code>
*	intersection	<code>int8range(5,15) * int8range(10,20)</code>	<code>[10,15)</code>
-	difference	<code>int8range(5,15) - int8range(10,20)</code>	<code>[5,10)</code>

LOTS EASIER

```
WHERE employed_during @@ [2018-01-01,2019-01-01)
```

VS

```
WHERE employed_from < '2019-01-01'  
AND    '2018-01-01' < employed_til
```


TWO DIMENSIONS

Valid Time	Transaction Time
history of the thing	history of the database
application features	auditing, compliance
user can edit	immutable
maintained by your app	maintained by triggers
constraints matter	look Ma, no hands!
nothing	pg: <code>temporal_tables</code> , "A Tardis for Your ORM", <code>pg_audit_log</code>
nothing	Rails: <code>papertrail</code> , <code>audited</code> , <code>chronomodel</code>


TERMINOLOGY

Snodgrass	valid time	transaction time
Fowler	actual time	record time
Date/Darwen/Lorentzos	stated time	logged time
Johnston	effective time/ state time	assertion time
SQL:2011	application time	system time

NON-UNIQUE PKS

products			
id	name	price	valid_at
1 1	shoe shoe	\$5 \$7	[Jan 2018, Jan 2019) [Jan 2019,)
2	snow	\$2	[Jan 2016,)
3 3	sail sail	\$8 \$9	[Jan 2016,) [Jan 2017, Jan 2018)

ANOTHER BOOK



EXCLUSION CONSTRAINTS

```
ALTER TABLE products  
ADD CONSTRAINT pk_products  
EXCLUDE USING gist  
(id WITH =, valid_at WITH &&);
```

EXCLUSION CONSTRAINTS

```
ALTER TABLE products  
ADD CONSTRAINT pk_products  
EXCLUDE  
(id WITH =);
```

EXCLUSION CONSTRAINTS

```
ALTER TABLE products  
ADD CONSTRAINT pk_products  
EXCLUDE USING gist  
(id WITH =, valid_at WITH &&);
```

EXCLUSION CONSTRAINTS

```
CHECK (  
  NOT EXISTS (  
    SELECT 1  
    FROM   products t1  
    WHERE  1 < (  
      SELECT COUNT(id)  
      FROM   products t2  
      WHERE  t1.id = t2.id  
      AND    t1.valid_at && t2.valid_at))  
  AND NOT EXISTS (  
    SELECT 1  
    FROM   products t3  
    WHERE  p3.id IS NULL)  
)
```


FOREIGN KEYS

products		
id	name	price
1	shoe	\$5
2	snow	\$2



variants		
id	product_id	size
1 2	1 1	5 8
3	2	5
4	3	1

TEMPORAL FOREIGN KEYS

products			
id	name	price	valid_at
1	shoe	\$5	[Jan 2018, Jan 2019)
1	shoe	\$7	[Jan 2019,)
2	snow	\$2	[Jan 2016,)



variants			
id	product_id	size	valid_at
1	1	5	[Jan 2018, Mar 2018)
2	1	8	[Jan 2018, Jan 2020)
3	2	5	[Jan 2014, Jan 2015)

TEMPORAL FOREIGN KEYS

```
CHECK (  
  NOT EXISTS (  
    SELECT 1  
    FROM    variants AS v  
    -- There was a p when v started:  
    WHERE NOT EXISTS (  
      SELECT 1  
      FROM    products AS p  
      WHERE   v.product_id = p.id  
      AND     coalesce(lower(p.valid_at), '-infinity')  
              <= coalesce(lower(v.valid_at), '-infinity')  
      AND     coalesce(lower(v.valid_at), '-infinity')  
              <= coalesce(upper(p.valid_at), 'infinity'))  
    -- ...  
  )  
)
```

TEMPORAL FOREIGN KEYS

```
-- ...
-- There was a p when v ended:
OR NOT EXISTS (
  SELECT 1
  FROM    products AS p
  WHERE   v.product_id = p.id
  AND     coalesce(lower(p.valid_at), '-infinity')
          < coalesce(upper(v.valid_at), 'infinity')
  AND     coalesce(upper(v.valid_at), 'infinity')
          <= coalesce(upper(p.valid_at), 'infinity'))
-- ...
```

TEMPORAL FOREIGN KEYS

```
-- ...
-- There are no gaps in p throughout v:
OR EXISTS (
  SELECT 1
  FROM    products AS p
  WHERE   v.product_id = p.id
  AND     coalesce(lower(v.valid_at), '-infinity')
          < coalesce(upper(p.valid_at), 'infinity')
  AND     coalesce(upper(p.valid_at), 'infinity')
          < coalesce(upper(v.valid_at), 'infinity')
-- ...
```

TEMPORAL FOREIGN KEYS

```
-- ...  
AND NOT EXISTS (  
    SELECT 1  
    FROM    products AS p2  
    WHERE   p2.id = p.id  
    AND     coalesce(lower(p2.valid_at), '-infinity')  
            <= coalesce(upper(p.valid_at), 'infinity')  
    AND     coalesce(upper(p.valid_at), 'infinity')  
            < coalesce(upper(p2.valid_at), 'infinity'))))
```

QUERIES

snapshot ("current")	at a given moment	returns a traditional table (removes <code>valid_at</code>)	<code>WHERE valid_at @> t</code>
sequenced	across time	returns a temporal table (preserves <code>valid_at</code>)	nothing, or <code>WHERE valid_at && r</code>
non- sequenced	time is just another column	returns ???	

JOINS

offers		
house_id	price	valid_at
1	\$100	[Feb 11, Feb 14)
1	\$150	[Feb 14, Feb 17)
1	\$100	[Feb 17, Feb 22)



reservations		
house_id	customer_id	valid_at
1	1	[Feb 13, Feb 15)
1	2	[Feb 16, Feb 19)

JOINS



offers		
house_id	price	valid_at
1	\$100	[Feb 11, Feb 14)
1	\$150	[Feb 14, Feb 17)
1	\$100	[Feb 17, Feb 22)



reservations		
house_id	customer_id	valid_at
1	1	[Feb 13, Feb 15)
1	2	[Feb 16, Feb 19)



customer_id	price	valid_at
1	\$100	[Feb 13, Feb 14)
1	\$150	[Feb 14, Feb 15)
2	\$150	[Feb 16, Feb 17)
2	\$100	[Feb 17, Feb 19)

JOINS





11 12 13 14 15 16 17 18 19 20 21 22
Feb

JOINS




11 12 13 14 15 16 Feb 17 18 19 20 21 22




11 12 13 14 15 16 Feb 17 18 19 20 21 22


TEMPORAL INSERT




TEMPORAL INSERT




TEMPORAL INSERT



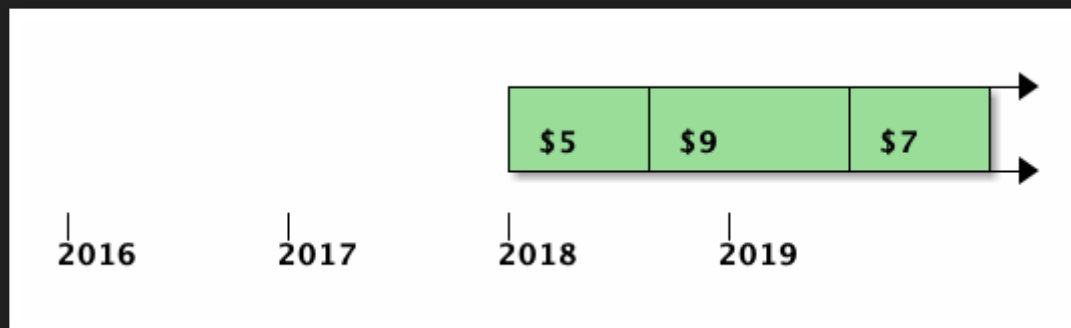

TEMPORAL INSERT



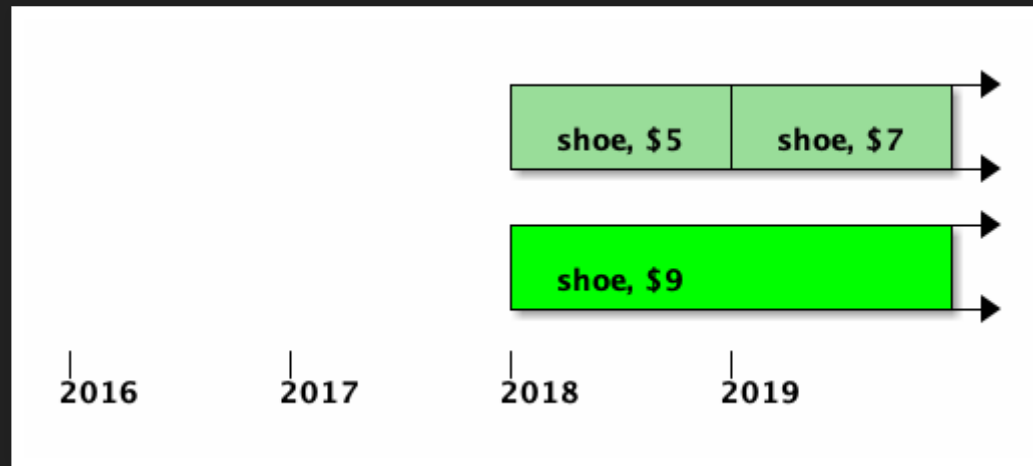
TEMPORAL UPDATE




TEMPORAL UPDATE



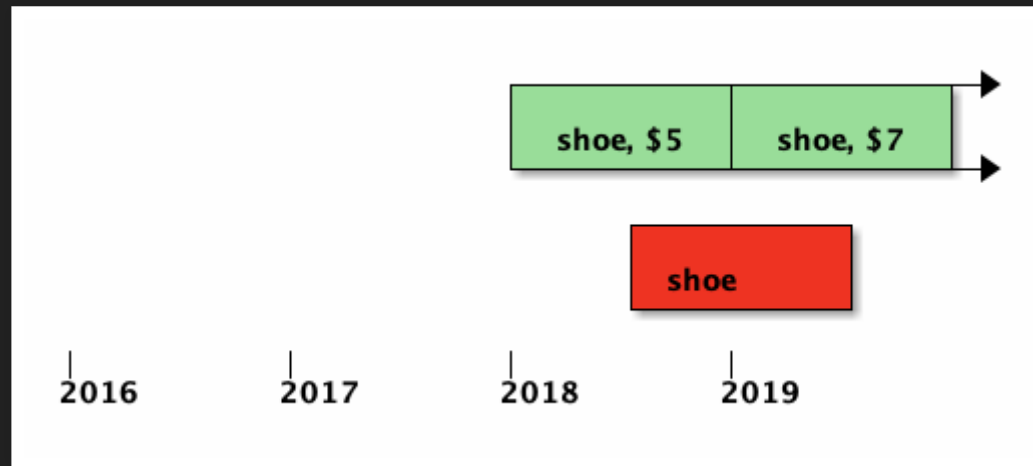
TEMPORAL UPDATE



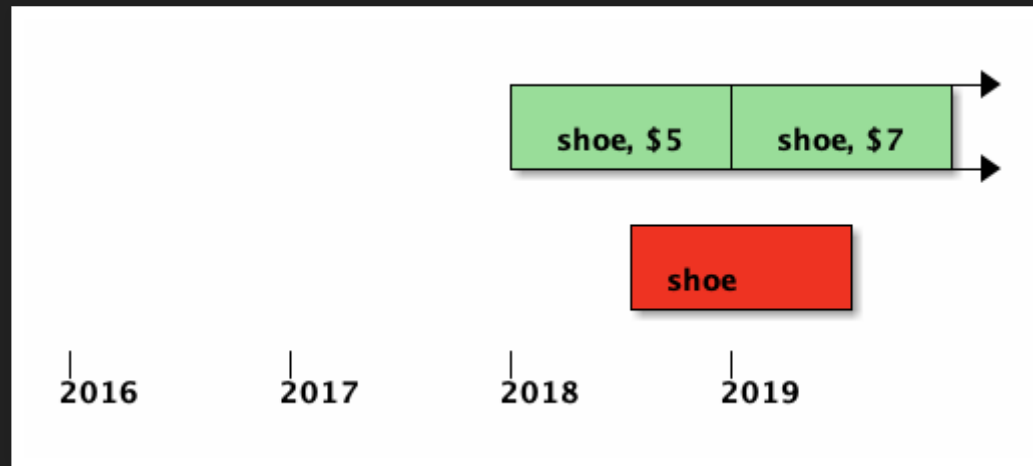
TEMPORAL UPDATE



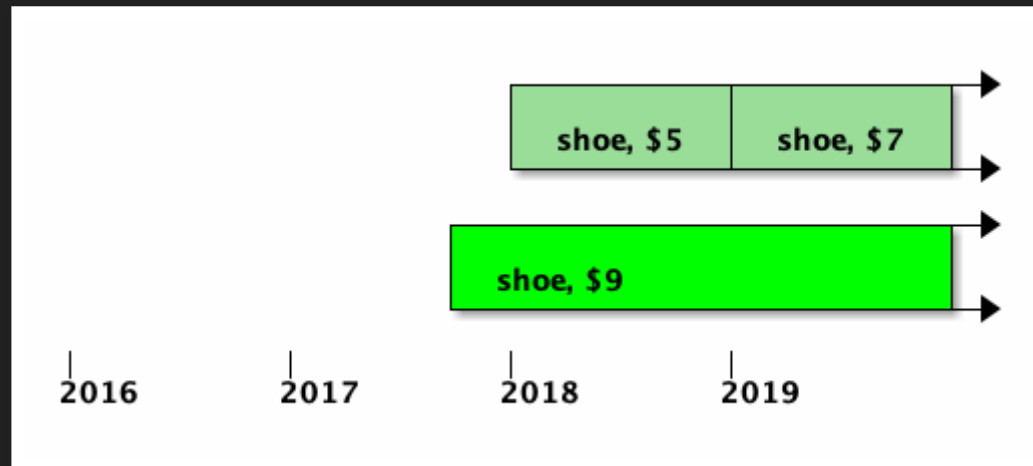
TEMPORAL DELETE




TEMPORAL DELETE



TEMPORAL UPSERT



TEMPORAL UPSERT



SQL:2011

RANGE

```
CREATE TABLE products (  
  id          integer,  
  valid_at   tstzrange,  
  
  name       text,  
  price      decimal(10,2),  
  
  CONSTRAINT pk_products  
    EXCLUDE  
    (id WITH =, valid_at WITH &&)  
);
```

PERIOD

```
CREATE TABLE products (  
  id          integer,  
  valid_from  timestampz NOT NULL,  
  valid_til   timestampz NOT NULL,  
  
  name        text,  
  price       decimal(10,2),  
  
  PERIOD FOR valid_at  
    (valid_from, valid_til),  
  CONSTRAINT pk_products  
    PRIMARY KEY  
    (id, valid_at WITHOUT OVERLAPS)  
);
```

PERIODS

```
SELECT * FROM t;
```

RANGES

```
CREATE TABLE products (  
  id          integer,  
  valid_at    tstzrange,  
  
  name        text,  
  price       decimal(10,2),  
  
  CONSTRAINT pk_products  
    PRIMARY KEY  
    (id, valid_at WITHOUT OVERLAPS)  
);
```

DML

```
INSERT INTO products (id, price, valid_at)
VALUES (1, 5, tstzrange(...));
```

```
UPDATE products
FOR PORTION OF valid_at FROM t1 TO t2
SET price = 4
WHERE id = 1;
```

```
DELETE FROM products
FOR PORTION OF valid_at FROM t1 TO t2
WHERE id = 1;
```

SYSTEM TIME

```
CREATE TABLE products (  
  id          integer,  
  sys_from    timestamp GENERATED ALWAYS AS ROW START,  
  sys_til     timestamp GENERATED ALWAYS AS ROW END,  
  
  name        text,  
  price       decimal(10,2),  
  
  PERIOD FOR SYSTEM_TIME  
    (sys_from, sys_til)  
) WITH SYSTEM VERSIONING;
```

SYSTEM TIME

```
SELECT *  
FROM   products  
FOR SYSTEM_TIME AS OF t;
```

```
SELECT *  
FROM   products  
FOR SYSTEM_TIME FROM t1 TO t2;
```

MARIADB

```
CREATE TABLE products (  
  id      integer,  
  name    text,  
  price   decimal(10,2)  
) WITH SYSTEM VERSIONING;  
  
SELECT *, row_start, row_end FROM products;
```


ORACLE

```
ALTER DATABASE ADD SUPPLEMENTAL LOG DATA;
```

```
SELECT * FROM t AS OF TIMESTAMP t;
```

ORACLE

```
CREATE TABLE products (  
  id          int,  
  valid_from  timestamp,  
  valid_til   timestamp,  
  
  name        text,  
  price       decimal(10,2),  
  
  PERIOD FOR valid_at (valid_from, valid_til)  
);  
  
SELECT * FROM products  
AS OF PERIOD FOR valid_at t;  
  
SELECT * FROM products  
VERSIONS PERIOD FOR valid_at BETWEEN t1 AND t2;
```

MS SQL SERVER

```
CREATE TABLE dbo.Products (  
    Id          integer PRIMARY KEY CLUSTERED,  
    SysFrom datetime2(2) GENERATED ALWAYS AS ROW START,  
    SysTil  datetime2(2) GENERATED ALWAYS AS ROW END,  
    Name      text,  
    Price     decimal(10, 2),  
    PERIOD FOR SYSTEM_TIME (SysFrom, SysTil)  
) WITH (  
    SYSTEM VERSIONING = ON  
        (HISTORY TABLE = dbo.ProductsHistory)  
);
```

IBM DB2

```
CREATE TABLE products (  
    id          int,  
    sys_from    timestamp(12) GENERATED ALWAYS AS ROW BEGIN,  
    sys_til     timestamp(12) GENERATED ALWAYS AS ROW END,  
    tx_id       timestamp(12) GENERATED ALWAYS AS TRANSACTION START  
  
    name        text,  
    price       decimal(10,2),  
  
    PERIOD SYSTEM_TIME (sys_from, sys_til)  
);
```

IBM DB2

```
CREATE TABLE products (  
    id          int,  
    valid_from  timestamp(12) NOT NULL,  
    valid_til   timestamp(12) NOT NULL,  
  
    name  text,  
    price decimal(10,2),  
  
    PERIOD BUSINESS_TIME (valid_from, valid_til)  
);
```

POSTGRES?

```
CREATE TABLE products (  
  id          integer,  
  claimed_at  tstzrange GENERATED ALWAYS AS ROW RANGE,  
  name        text,  
  price       decimal(10,2)  
) WITH SYSTEM VERSIONING (claimed_at);
```

THANKS!

ME

- <https://github.com/pjungwir/temporal-databases-postgres-talk>
- <https://illuminatedcomputing.com/posts/2017/12/temporal-databases-bibliography/>

RESEARCH

- <https://www2.cs.arizona.edu/~rts/publications.html>
- <http://citeseerx.ist.psu.edu/viewdoc/download;jsessionid=F78723B857463955C76E540DCAB8FDF5?doi=10.1.1.116.7598&rep=rep1&type=pdf>
- <https://files.ifi.uzh.ch/boehlen/Papers/modf174-dignoes.pdf>
- http://www.zora.uzh.ch/id/eprint/130374/1/Extending_the_kernel.pdf

SQL:2011

- <https://www.wiscorp.com/SQLStandards.html>
- <https://sigmodrecord.org/publications/sigmodRecord/1209/pdfs/07.industry.kulkarni.pdf>

OTHER VENDORS

- <https://mariadb.com/kb/en/library/system-versioned-tables/>
- https://docs.oracle.com/database/121/ADFNS/adfns_flashback.htm#ADFNS610
- https://docs.oracle.com/database/121/ADFNS/adfns_design.htm#ADFNS967
- <https://docs.microsoft.com/en-us/sql/relational-databases/tables/temporal-tables?view=sql-server-2017>
- https://www.ibm.com/support/knowledgecenter/en/SSEPGG_10.1.0/com.ibm.db2.luw.admin.dbobj.doc/doc/t0058926.html

PATCHES

- <https://www.postgresql-archive.org/PROPOSAL-Temporal-query-processing-with-range-types-tt5913058.html>
- <https://www.postgresql-archive.org/SQL-2011-PERIODS-vs-Postgres-Ranges-tt6055264.html>

TOOLS

- https://github.com/arkhipov/temporal_tables
- <https://www.youtube.com/watch?v=TRgni5q0YM8>
- <https://github.com/ifad/chronomodel>